The Architect of Chemical Space: A Comprehensive Analysis of John Irwin's Contributions to Structural Biology and Computational Drug Discovery

Executive Summary

The trajectory of modern pharmaceutical science has been defined by a transition from serendipitous observation to rational, structure-guided design. At the heart of this transformation lies the digitalization of chemical space—the ability to represent, organize, and screen billions of molecules in silico before a single gram is synthesized in the laboratory. Dr. John J. Irwin, an Adjunct Professor of Pharmaceutical Chemistry at the University of California, San Francisco (UCSF), stands as a central figure in this revolution. Through a career spanning over three decades, Irwin has engineered the foundational infrastructure upon which much of contemporary academic and industrial drug discovery relies.

This report provides an exhaustive examination of Irwin's scientific contributions, specifically tailored for an advanced biochemical audience. It explores his development of the ZINC database, which democratized access to virtual screening libraries; his work on the Similarity Ensemble Approach (SEA), which redefined our understanding of polypharmacology; and his critical investigations into colloidal aggregation, which identified and mitigated a primary source of false positives in high-throughput screening. Furthermore, it details his recent ventures into ultra-large library docking and the founding of Deep Apple Therapeutics, efforts that integrate cryo-electron microscopy and deep learning to accelerate the discovery of ligands for G-protein-coupled receptors (GPCRs). By analyzing Irwin's methodologies, collaborations, and the theoretical underpinnings of his software tools, this document illustrates how computational chemoinformatics has become an indispensable pillar of structural biology.

1. Introduction: The Digital Transformation of Biochemistry

1.1 The Paradigm Shift in Ligand Discovery

For much of the 20th century, the discovery of new therapeutic agents was largely an empirical process. Chemists synthesized analogs of natural products or known active compounds, and biologists tested them in phenotypic assays. While successful, this approach was labor-intensive and limited by the physical availability of chemical matter. The "chemical

space"—the theoretical set of all possible stable organic molecules below 500 Daltons—is estimated to contain between 10^{60} and 10^{200} entities. Physical screening, even with high-throughput robotics, can typically assess only a few million compounds, a vanishingly small fraction of the possibilities.

The advent of structural biology, particularly X-ray crystallography and later cryo-electron microscopy (cryo-EM), provided a new map for discovery. By resolving the three-dimensional atomic structures of protein targets, scientists could theoretically design molecules to fit into binding pockets like keys into locks. However, this Structure-Based Drug Design (SBDD) required a bridge between the physical world of protein structures and the digital world of chemical possibilities. Building this bridge required not just computational power, but sophisticated data architecture capable of handling the stereochemical and physical complexity of millions, and eventually billions, of molecules.

1.2 Defining the Subject

Dr. John J. Irwin emerged as the architect of this bridge. His work does not merely involve the creation of databases but encompasses the rigorous physical simulation of molecular properties necessary for accurate virtual screening. It is important to distinguish the subject of this report, the computational chemist John Irwin, from the noted sociologist and criminologist John Irwin (1929–2010), who was also associated with the California university system and founded Project Rebound.[1] The John Irwin discussed herein is a distinct figure, trained in chemical crystallography and currently active in the Department of Pharmaceutical Chemistry at UCSF, often in close collaboration with Professor Brian Shoichet.[3]

Irwin's work addresses the fundamental "data problem" in biochemistry. As structural data proliferated via the Protein Data Bank (PDB) and chemical synthesis capabilities expanded, the gap between what could be simulated and what could be physically tested widened. Irwin's contributions—ZINC, DOCK Blaster, DUD-E, SEA, and the Aggregation Advisor—systematically closed this gap, providing the tools that allow a biology student with a laptop to perform screening campaigns that once required the resources of a major pharmaceutical company.[5]

2. Foundations in Structural Science and Crystallography

2.1 The Dunitz Era: Rigor in Three Dimensions

John Irwin's scientific philosophy was forged in the exacting discipline of chemical crystallography. He earned his Ph.D. in 1991 from ETH Zurich (Swiss Federal Institute of Technology), working under the mentorship of Jack Dunitz.[3] Jack Dunitz was a towering figure in structural chemistry, renowned for his work on crystal packing, stereochemistry, and the

analysis of molecular motion in the solid state.

This training is critical to understanding Irwin's later work in chemoinformatics. In the early days of computational chemistry, many databases treated molecules as two-dimensional graphs—flat representations of atoms and bonds. However, biological recognition is an inherently three-dimensional event. A ligand must adopt a specific conformation to bind to a protein, and the energy penalty for adopting that conformation contributes to the binding affinity. Irwin's background in crystallography instilled a deep appreciation for the physical reality of molecules—their bond lengths, angles, torsion strains, and packing forces. This crystallographic rigor would become a defining feature of the ZINC database, distinguishing it from competitors by ensuring that every molecule was represented in a biologically relevant, 3D form ready for physical simulation.[3]

2.2 The Cambridge Years: Software for Structure Solution

Following his doctoral work, Irwin transitioned into software development for structural biology, a move that placed him at the center of the field's computational evolution. He joined the laboratory of Gerard Bricogne at the Medical Research Council (MRC) Laboratory of Molecular Biology in Cambridge, UK.[3] The MRC LMB is historically significant as the birthplace of molecular biology, associated with figures like Crick, Watson, and Perutz.

In Bricogne's group, Irwin contributed to the development of two seminal software packages: **SHARP** and **BUSTER**.[3]

- **SHARP (Statistical Heavy-Atom Refinement and Phasing):** Solving a crystal structure requires determining the "phase" of the diffracted X-rays, information that is lost during the experiment. SHARP used maximum likelihood methods to locate heavy atoms (used as markers) and estimate phases with high statistical rigor. This revolutionized de novo structure determination.
- **BUSTER:** This software focused on the refinement of the structural model, optimizing the fit of the atomic coordinates to the electron density map.

Working on these tools required Irwin to master the complex mathematics of diffraction and the optimization algorithms necessary to handle noisy experimental data. It also reinforced the concept that computational models are only as good as the data and the physical principles underlying them.

2.3 Data Curation at EMBL-EBI

Before moving to the United States, Irwin spent approximately 18 months as a staff scientist at the European Molecular Biology Laboratory - European Bioinformatics Institute (EMBL-EBI), also located in Cambridge.[3] His primary focus was the "PDB cleanup project."

The Protein Data Bank (PDB) is the global archive for 3D structural data. In the late 1990s, the archive suffered from inconsistencies in data formats, nomenclature, and annotation. Irwin's

work involved developing automated pipelines to standardize and "clean" these files, ensuring they were machine-readable and chemically consistent. This experience was a masterclass in large-scale data curation. It exposed him to the heterogeneity of biological data and the necessity of robust, automated systems to manage it—skills that would be directly applicable to the curation of chemical catalogs in the ZINC project.[3]

2.4 The Shoichet Collaboration and the Move to UCSF

In 2000, Irwin moved to Northwestern University Medical School to join the lab of Brian Shoichet.[3] This marked the beginning of a prolific professional partnership that has spanned over two decades. When the Shoichet Lab moved to UCSF in the mid-2000s, Irwin moved with it, taking up a position as Adjunct Associate Professor in the Department of Pharmaceutical Chemistry.[3]

At UCSF, Irwin found himself in an environment steeped in the history of structure-based drug design. UCSF was the home of Tack Kuntz, who developed the original DOCK algorithm in 1982. The environment was highly collaborative, merging experimental pharmacology (with colleagues like Bryan Roth) and computational theory. Irwin's role evolved into that of a "technological enabler"—the person who built the infrastructure that allowed theoretical docking methods to be applied to real-world drug discovery campaigns.[7]

3. The Architecture of Virtual Screening: UCSF DOCK and Benchmarking

3.1 The Principles of Molecular Docking

To appreciate Irwin's contribution, one must understand the computational problem of docking. Molecular docking aims to predict the preferred orientation of a ligand to a protein (pose prediction) and the strength of the binding (scoring). The Shoichet Lab, and Irwin specifically, utilize and develop **UCSF DOCK** (currently versions 3.7/3.8).[3]

Unlike many commercial docking programs that use stochastic methods (like genetic algorithms) to explore ligand flexibility during the docking process, UCSF DOCK 3.7 typically pre-calculates ligand conformations. This places a massive burden on the database preparation—the ligands must be generated in thousands of flexible shapes *before* the docking run begins. This architectural choice necessitates a highly structured database like ZINC, which can store and serve these pre-computed conformations efficiently.

3.2 DOCK Blaster: Democratizing the Algorithm

While DOCK is powerful, it is chemically and computationally complex, operating primarily via command-line interfaces on Linux clusters. This creates a barrier to entry for biologists who may have a crystal structure but lack computational expertise. To bridge this divide, Irwin

developed **DOCK Blaster**, a web-based public access service.[3]

DOCK Blaster allows a user to upload a PDB code (the identifier for their protein structure). The server then:

1. Analyzes the protein structure to identify the binding site (often using the location of a co-crystallized ligand).
2. Prepares the receptor (calculating molecular surfaces and energy grids).
3. Automatically selects a subset of the ZINC database to screen.
4. Runs the docking simulation on a cluster.
5. Returns a ranked list of compounds that the biologist can purchase and test.

This tool was revolutionary for introductory biochemistry contexts and small labs, as it effectively provided a "computational chemist in a box," allowing researchers to generate testable hypotheses about ligand binding without hiring a specialist.[10]

3.3 The Crisis of Evaluation: DUD and DUD-E

A recurrent problem in computational chemistry is validation: How do we know a docking program actually works? Early benchmarks often used a set of "active" molecules (known binders) and a set of "decoy" molecules (assumed non-binders) chosen at random from a database.

Irwin and his colleagues identified a critical flaw in this approach: **Artificial Enrichment via Property Bias**. If the active molecules for a target were large and greasy (hydrophobic), and the random decoys were small and polar, a docking program could "succeed" simply by scoring large/greasy molecules better, without actually calculating the specific fit to the protein pocket. The program was getting the right answer for the wrong reason.[12]

To solve this, Irwin's team developed the **Directory of Useful Decoys (DUD)** and later the **Directory of Useful Decoys, Enhanced (DUD-E)**.[3]

- **Methodology:** For every active ligand, DUD-E selects 50 decoy molecules from the ZINC database. These decoys are "property-matched"—they have similar molecular weight, calculated LogP (lipophilicity), and number of rotatable bonds as the active ligand. However, they have dissimilar topological structures (2D fingerprints) to ensure they are likely non-binders.
- **Significance:** By matching physical properties, DUD-E forces the docking algorithm to discriminate based on *molecular recognition* (the precise 3D fit and specific interactions) rather than gross physical properties.
- **Impact:** DUD-E is now the industry standard for benchmarking. Any new docking algorithm or scoring function developed today must demonstrate its performance on DUD-E to be taken seriously by the community. It represents a fundamental contribution to the rigorous scientific method within the field.[12]

4. The ZINC Database: A Chemist’s Universe

4.1 The Genesis of ZINC

Prior to 2005, finding purchasable compounds for screening was a fragmented and expensive process. Databases like the Available Chemicals Directory (ACD) existed but were costly and often lacked the 3D data required for docking. Irwin and Shoichet launched ZINC (a recursive acronym for **Z**INC **I**s **N**ot **C**ommercial) to solve this.[6]

The initial goal was simple: creating a free, curated collection of commercially available compounds, formatted specifically for virtual screening. The 2005 launch paper in the *Journal of Chemical Information and Modeling* introduced a library of approximately 700,000 compounds.[13] This was the "In-Stock" era, where the database reflected molecules that were physically sitting in bottles in vendor warehouses.

4.2 The Philosophy of "Ready-to-Dock"

ZINC was not just a list of names; it was a physical simulation of the molecules. Irwin implemented rigorous protocols to prepare these molecules, addressing chemical nuances that often trip up simulations:

- **Protonation and Tautomers:** Biological pH (typically 7.4) dictates the charge state of a molecule. A carboxylic acid is deprotonated (negative), and an amine is protonated (positive). ZINC uses algorithms to predict these states, generating multiple forms if the pKa is near physiological pH. This is critical because an incorrect charge state leads to massive errors in electrostatic scoring during docking.[6]
- **3D Conformation:** ZINC generates 3D coordinates (using tools like OpenEye's Omega or Corina), saving users the computational cost of generating them.
- **Subsetting:** To improve efficiency, Irwin organized ZINC into subsets based on physicochemical properties, such as "Fragment-like" (MW < 250), "Lead-like" (MW 250-350), and "Drug-like" (MW 350-500, Lipinski compliant). This allowed users to screen only the chemical space relevant to their specific project goals.[5]

4.3 The Make-on-Demand Revolution (ZINC15 to ZINC20)

The most profound shift in Irwin’s work occurred with the advent of "Make-on-Demand" libraries. Around 2015, chemical vendors, led by **Enamine Ltd.** in Kyiv, Ukraine, changed their business model. Instead of synthesizing compounds and storing them, they listed molecules they *could* synthesize using verified chemical reactions and available building blocks.[8]

This created the **REAL (Readily Accessible)** database. These molecules did not exist physically but could be manufactured with a success rate exceeding 80% within a few weeks.

- **The Scale Up:** This shifted the available chemical space from millions (in-stock) to billions (virtual/tangible). ZINC15 contained ~120 million compounds. ZINC20 expanded

this to over 1.4 billion.[6]

- **The Implications:** This was not just "more of the same." Analysis by Irwin and colleagues showed that these make-on-demand libraries contained fundamentally new shapes (scaffolds) and topologies that were absent from historic libraries. This expansion meant that for difficult protein targets, where no known binders existed, the probability of finding a match increased simply because the diversity of shapes being tested had increased orders of magnitude.[16]

4.4 ZINC22 and Federated Data

As the library grew to over 37 billion compounds in **ZINC22**, traditional database architectures failed. A single SQL database could not efficiently index or search this volume of data. Irwin engineered a federated approach, sharding the database into smaller, static files based on chemical properties (e.g., Heavy Atom Count and LogP). This allowed the database to be hosted on cloud storage (like Amazon S3 via the AWS Open Data Sponsorship Program) and searched using massive parallelization.[17]

Table 1: The Evolution of the ZINC Database

Version	Year	Approx. Size	Key Innovation	Scientific Context
ZINC (Original)	2005	~700,000	Free access, 3D formats	Introduction of "Ready-to-Dock" concept.[13]
ZINC12	2012	~20 Million	Subsets (Lead-like vs Drug-like)	Refinement of screening for specific physicochemical profiles.[10]
ZINC15	2015	~120 Million	Integration of "Tangible" libraries	The beginning of the make-on-demand transition.[6]
ZINC20	2020	~1.4 Billion	Ultra-Large Library support	Full integration with Enamine REAL space.[16]

ZINC22	2022	~37 Billion	Federated Architecture	Handling "inexhaustible" chemical space via cloud computing.[18]

5. The Quality Control of Chemical Space: Aggregation and Artifacts

5.1 The False Positive Crisis

In High-Throughput Screening (HTS), researchers often test hundreds of thousands of compounds against an enzyme target. It is common to find "hits"—compounds that inhibit the enzyme. However, a significant percentage of these hits are artifacts; they inhibit the enzyme not by binding to a specific pocket, but by non-specific mechanisms. Pursuing these artifacts wastes millions of dollars and years of research effort.

5.2 The Biophysics of Colloidal Aggregation

Irwin and Shoichet identified **colloidal aggregation** as the dominant mechanism for these false positives. Many organic molecules, at micromolar concentrations, spontaneously self-assemble into microscopic, liquid-like droplets (colloids). These colloids inhibit enzymes by:

1. **Sequestration:** The enzyme adsorbs to the surface of the colloid, removing it from the solution.
2. **Denaturation:** The interaction with the colloid surface can partially unfold the protein.

This inhibition is "promiscuous"—an aggregator will inhibit Beta-lactamase, Kinases, and GPCRs indiscriminately.[20]

5.3 The Aggregation Advisor

Irwin turned this physical insight into a predictive tool. He curated a database of known aggregators (molecules proven to aggregate in experimental assays). He then developed the **Aggregation Advisor**, a tool integrated into ZINC.[21]

- **Mechanism:** When a user queries a molecule, the Advisor calculates its Tanimoto similarity (Tc) to the known aggregators and analyzes its calculated LogP (lipophilicity).
- **Prediction:** The tool flags molecules that look like known aggregators. It advises the user to perform a simple experimental control: add a non-ionic detergent (like Triton X-100) to the assay. Detergents disrupt colloids. If the inhibition disappears in the presence of detergent, the compound is an aggregator (an artifact). If inhibition persists, it is likely a

true binder.[21]

This work fundamentally changed how medicinal chemists evaluate HTS hits. "Aggregation" became a standard checklist item for validating new inhibitors, and Irwin's tool provided the means to check it rapidly.

6. Beyond the Lock and Key: Systems Pharmacology and SEA

6.1 The Polypharmacology Paradigm

The classical view of drug action is "one drug, one target." However, effective drugs often bind to multiple targets (polypharmacology), which can drive efficacy (e.g., in psychiatric drugs) or toxicity (side effects). Conversely, molecules that look nothing alike can bind to the same target.

6.2 The Similarity Ensemble Approach (SEA)

Working with Michael Keiser (a former student in the lab), Irwin developed the **Similarity Ensemble Approach (SEA)**.[3] SEA applies the logic of bioinformatics (sequence alignment) to chemoinformatics.

- **The Logic:** Just as BLAST compares protein sequences to find evolutionary relationships, SEA compares sets of ligands. If Protein A binds a set of ligands, and Protein B binds a set of ligands that are chemically similar to Protein A's ligands, then Protein A and Protein B are related by their chemistry, even if their amino acid sequences are totally different.
- **The E-Value:** SEA calculates an Expectation Value (E-value), a statistical measure of how likely the observed similarity between ligand sets would occur by random chance. A strong E-value suggests a functional link between targets.[26]

6.3 Applications: De-orphaning and Repurposing

Irwin and collaborators have used SEA for:

- **Drug Repurposing:** Predicting that known drugs bind to off-targets. For example, predicting that certain sedatives bind to dopamine receptors, explaining their side effect profiles.[27]
- **De-orphaning:** Predicting ligands for "orphan" GPCRs (receptors with no known endogenous ligand) by linking them to chemically related receptors.[28]
- **Toxicity Prediction:** Screening drug candidates against panels of targets (like the hERG channel) to predict adverse events before clinical trials.

7. The Frontier: Ultra-Large Library Docking

7.1 Testing the "Numbers Game"

Around 2019, with the integration of Enamine's billion-scale libraries into ZINC, Irwin and Shoichet asked a fundamental question: Does screening *more* molecules actually yield *better* results, or does it just produce more mediocre hits?

7.2 The Nature Papers (2019, 2020)

They answered this in a series of high-profile publications.[8]

- **The AmpC and D4 Study:** They docked 138 million compounds against AmpC Beta-lactamase and the D4 Dopamine receptor.
  - **Results:** They found hits with low nanomolar potency directly from docking.
  - **Novelty:** Critically, these hits were new chemotypes—they did not resemble known dopamine agonists. This proved that expanding the library allowed access to "islands" of chemical space that had never been explored.[8]
- **The Melatonin Study:** In 2020, they expanded this to the MT1 and MT2 melatonin receptors, screening hundreds of millions of compounds to find selective agonists that could modulate circadian rhythms in mice.[14]

7.3 Collaboration with Synthesis Experts

These projects were not purely computational. They involved tight collaboration with synthetic chemists like **Jonathan Ellman** (Yale) and pharmacologists like **Bryan Roth** (UNC Chapel Hill).[30] For example, in the study of the 5-HT2A receptor (the target of psychedelics), they used "bespoke" libraries—virtual libraries built around specific chemical reactions available in Ellman's lab—to find non-hallucinogenic antidepressants.[32]

8. Commercial Translation: Deep Apple Therapeutics

8.1 From Academic Code to Industrial Application

While Irwin's tools are free for academics, the commercial potential of finding new drugs using these methods is immense. Irwin has facilitated this transition through two primary commercial vehicles.

8.2 Blue Dolphin Lead Discovery

Co-founded with Brian Shoichet, **Blue Dolphin Lead Discovery** operates as a Contract Research Organization (CRO).[18]

- **Model:** It allows pharmaceutical companies to pay for "docking as a service." Companies utilize the Shoichet Lab's advanced protocols and the ZINC database without the IP complications of an academic collaboration.
- **Conflict of Interest Management:** The company is carefully structured to ensure

projects do not overlap with university research, maintaining scientific integrity.[35]

8.3 Deep Apple Therapeutics

In December 2023, Irwin, Shoichet, and **Georgios Skiniotis** (a structural biologist at Stanford specializing in Cryo-EM) launched **Deep Apple Therapeutics** with a $52 million Series A investment from Apple Tree Partners.[36]

- **The Convergence:** Deep Apple represents the industrial maturation of Irwin's life work. It combines:
  1. **Cryo-EM:** To generate high-resolution structures of GPCRs in active states (Skiniotis's contribution).
  2. **Ultra-Large Library Docking:** Using ZINC-scale libraries to screen billions of compounds (Irwin and Shoichet's contribution).
  3. **Deep Learning:** To accelerate the scoring of these massive libraries, allowing the screening of billions of compounds in hours.
- **Target:** The company focuses on GPCRs for metabolic and endocrine disorders, aiming to find oral small molecules that can replace peptide therapies.[36]

9. Conclusion

Dr. John J. Irwin has fundamentally altered the landscape of structural biology. Through the ZINC database, he moved the field from managing static lists of chemicals to navigating a dynamic, billion-scale universe of "tangible" molecules. His development of DOCK Blaster and SEA democratized high-performance computing, placing sophisticated drug discovery tools into the hands of biologists and students worldwide. Moreover, his rigorous investigation of artifacts through the Aggregation Advisor has imposed a necessary quality control standard on the field, preventing resources from being wasted on false positives.

For the biochemistry student, Irwin's work serves as a powerful demonstration of the synergy between disciplines. His career illustrates that the future of biology is not just in the wet lab, but in the server room, where the physics of molecular interactions meets the scale of big data. By building the infrastructure that connects protein structure to chemical supply, John Irwin has accelerated the pace at which humanity can discover new medicines.

Data Summary Tables

Table 2: Key Software and Tools Developed/Managed by John Irwin

Tool	Function	Key Scientific Concept
ZINC	Database of purchasable	3D conformation,

	compounds	protonation, "Make-on-Demand" scaling.[6]
DOCK Blaster	Web-based docking server	Democratization of HPC for non-specialists.[10]
SEA	Ligand-based target prediction	Polypharmacology, E-values for chemical similarity.[26]
DUD / DUD-E	Benchmarking datasets	Property-matched decoys to remove bias.[12]
Aggregation Advisor	Artifact prediction	Colloidal aggregation, false positive detection.[21]
Arthor / SmallWorld	Chemical search engines	Graph-based and similarity searching in billion-scale databases.[10]

Table 3: Major Collaborators and Roles

Collaborator	**Institution**	**Role in Irwin's Network**
Brian Shoichet	UCSF	Primary academic partner; docking theory and application.[3]
Bryan Roth	UNC Chapel Hill	Pharmacologist; validates GPCR hits in vitro/in vivo.[8]
Georgios Skiniotis	Stanford	Structural Biologist (Cryo-EM); co-founder of Deep Apple.[36]
Jonathan Ellman	Yale	Synthetic Chemist; bespoke library synthesis.[31]
Jack Dunitz	ETH Zurich	PhD Supervisor;

		foundational training in crystallography.[3]
Gerard Bricogne	MRC LMB	Postdoc Mentor; crystallographic software development.[3]

Works cited

1. Dr. John Irwin | Project Rebound - CSUSB, accessed January 26, 2026, https://www.csusb.edu/project-rebound/history/dr-john-irwin
2. John Irwin Dedication Video - 2012 - YouTube, accessed January 26, 2026, https://www.youtube.com/watch?v=gyt-FNnqj-s
3. Lab members - Irwin Lab - UCSF, accessed January 26, 2026, https://irwinlab.compbio.ucsf.edu/people
4. ZINC database - Wikipedia, accessed January 26, 2026, https://en.wikipedia.org/wiki/ZINC_database
5. ZINC database | PPTX - Slideshare, accessed January 26, 2026, https://www.slideshare.net/slideshow/zinc-database/134891423
6. ZINC 15 – Ligand Discovery for Everyone - PMC - NIH, accessed January 26, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC4658288/
7. John Irwin received his Ph.D. for work with Jack Dunitz in chemical crystallography in 1991 at ETH Zurich. He worked at a startu, accessed January 26, 2026, http://www.bch.cuhk.edu.hk/croucher09/Introduction_speakers.pdf
8. Creating an inexhaustible database for drug molecules - UCSF School of Pharmacy, accessed January 26, 2026, https://pharmacy.ucsf.edu/news/2019/02/creating-inexhaustible-database-drug-molecules
9. Large Scale Docking in the Cloud - ChemRxiv, accessed January 26, 2026, https://chemrxiv.org/doi/full/10.26434/chemrxiv-2023-f182x
10. Irwin Lab - Home - UCSF, accessed January 26, 2026, https://irwinlab.compbio.ucsf.edu/
11. Speaker Biographies - Discovery On Target, accessed January 26, 2026, https://www.discoveryontarget.com/19/speaker-biographies
12. by Submitted in partial satisfaction of the requirements for degree of in in the GRADUATE DIVISION of the UNIVERSITY OF CALIFORN - eScholarship.org, accessed January 26, 2026, https://escholarship.org/content/qt2768w8wt/qt2768w8wt_noSplash_1bcd1fec4fdf86e45d9f9d0f6bcefe01.pdf
13. ZINC: A Free Tool to Discover Chemistry for Biology | Journal of Chemical Information and Modeling - ACS Publications, accessed January 26, 2026, https://pubs.acs.org/doi/10.1021/ci3001277
14. ZINC20—A Free Ultralarge-Scale Chemical Database for Ligand Discovery - ACS Publications, accessed January 26, 2026, https://pubs.acs.org/doi/10.1021/acs.jcim.0c00675

15. Virtual drug discovery tackles unknown universe of drugs - DTR - Drug Target Review, accessed January 26, 2026, https://www.drugtargetreview.com/news/39890/drug-candidates/
16. ZINC20 – A Free Ultra Large-Scale Chemical Database for Ligand Discovery - eScholarship.org, accessed January 26, 2026, https://escholarship.org/content/qt0x7507b3/qt0x7507b3.pdf
17. (PDF) ZINC-22 A Free Multi-Billion-Scale Database of Tangible Compounds for Ligand Discovery - ResearchGate, accessed January 26, 2026, https://www.researchgate.net/publication/368542032_ZINC-22A_Free_Multi-Billion-Scale_Database_of_Tangible_Compounds_for_Ligand_Discovery
18. ZINC-22 A Free Multi-Billion-Scale Database of Tangible Compounds for Ligand Discovery | Journal of Chemical Information and Modeling, accessed January 26, 2026, https://pubs.acs.org/doi/10.1021/acs.jcim.2c01253
19. ZINC Database - AWS Marketplace, accessed January 26, 2026, https://aws.amazon.com/marketplace/pp/prodview-cclrxhtx5xibk
20. Colloidal aggregation in drug discovery & drug formulation - Shoichet Lab, accessed January 26, 2026, https://bkslab.org/projects/aggregation
21. An Aggregation Advisor for Ligand Discovery - PubMed - NIH, accessed January 26, 2026, https://pubmed.ncbi.nlm.nih.gov/26295373/
22. Aggregation alert - Practical Fragments, accessed January 26, 2026, http://practicalfragments.blogspot.com/2015/10/aggregation-alert.html
23. An Aggregation Advisor for Ligand Discovery | Request PDF - ResearchGate, accessed January 26, 2026, https://www.researchgate.net/publication/281167262_An_Aggregation_Advisor_for_Ligand_Discovery
24. Modeling the expansion of virtual screening libraries - PMC - NIH, accessed January 26, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC10243288/
25. John Irwin, PhD - UCSF Profiles, accessed January 26, 2026, https://profiles.ucsf.edu/john.irwin
26. UCSF - UC San Francisco Electronic Theses and Dissertations - eScholarship.org, accessed January 26, 2026, https://escholarship.org/content/qt5vp5h8g4/qt5vp5h8g4.pdf
27. Method to predict polypharmacology developed in Shoichet Lab | School of Pharmacy, accessed January 26, 2026, https://pharmacy.ucsf.edu/news/2009/12/method-predict-polypharmacology-developed-shoichet-lab
28. IDG UNC meeting presentation by Dr. John Irwin (UCSF) - YouTube, accessed January 26, 2026, https://www.youtube.com/watch?v=a4gaqPqH7NM
29. Current Status of Computational Approaches for Small Molecule Drug Discovery | Journal of Medicinal Chemistry - ACS Publications, accessed January 26, 2026, https://pubs.acs.org/doi/10.1021/acs.jmedchem.4c02462
30. Mega Docking Library Poised to Speed Drug Discovery - National Institute of Mental Health, accessed January 26, 2026, https://www.nimh.nih.gov/news/science-updates/2019/mega-docking-library-poised-to-speed-drug-discovery

31. 2022 - Practical Fragments, accessed January 26, 2026, http://practicalfragments.blogspot.com/2022/
32. Creating new neuropsychiatric drug candidates — from a virtual library | Yale News, accessed January 26, 2026, https://news.yale.edu/2022/09/28/creating-new-neuropsychiatric-drug-candidates-virtual-library
33. John J Irwin Ph.D. Professor (Associate) at University of California, San Francisco - ResearchGate, accessed January 26, 2026, https://www.researchgate.net/profile/John-Irwin-4
34. Blue Dolphin Lead Discovery - 2025 Company Profile, Team & Competitors - Tracxn, accessed January 26, 2026, https://tracxn.com/d/companies/blue-dolphin-lead-discovery/__sMfPvcrZ-F024sC7fmh7WflEOgli1Azx-uEUSM4AxaY
35. Blue Dolphin Lead Discovery Virtual Screening Services - Founders, accessed January 26, 2026, http://bluedolphinleads.s3-website.us-east-2.amazonaws.com/founders.html
36. Apple Tree Partners Unveils Deep Apple Therapeutics to Dramatically Accelerate Drug Discovery with Structure-Based, Ultra-Large Library Virtual Screening and Deep Learning Models - PR Newswire, accessed January 26, 2026, https://www.prnewswire.com/news-releases/apple-tree-partners-unveils-deep-apple-therapeutics-to-dramatically-accelerate-drug-discovery-with-structure-based-ultra-large-library-virtual-screening-and-deep-learning-models-302014867.html
37. Deep Apple Therapeutics Launches With $52 Million In Funding, accessed January 26, 2026, https://pulse2.com/deep-apple-therapeutics-launches-with-52-million-in-funding/
38. Deep Apple launches | C&EN Global Enterprise - ACS Publications, accessed January 26, 2026, https://pubs.acs.org/doi/10.1021/cen-10201-buscon2